Comprehensive Speed Report

Abstract

The implementation of the multi-threaded server and its search algorithms is the result of extensive research and iterative testing. The primary goal was to ensure robust performance and scalability for files containing up to 1,000,000 rows while meeting strict execution time requirements: an average of 40 milliseconds for dynamic file searches (REREAD_ON_QUERY = True) and 0.5 milliseconds for cached file searches (REREAD_ON_QUERY = False).

Implementation Overview

Algorithm Selection and Optimization

Through detailed experimentation and performance profiling, a hybrid approach combining memory-mapped file access (mmap) and set-based caching was determined to provide the optimal balance of speed and resource utilization. These algorithms outperformed other tested methods, such as linear search, binary search, and B-Trees, in both dynamic and static search scenarios.

Key Results from Algorithm Testing:

Dynamic Searches (REREAD_ON_QUERY = True):

- Execution times averaged **18.82 ms** for file sizes up to 250,000 rows.
- Memory-mapped file access ensured minimal I/O overhead, allowing for consistent performance.

Cached Searches (REREAD_ON_QUERY = False):

- Execution times averaged **0.000088 ms**, well within the 0.5-millisecond target.
- The set structure provided constant-time lookups.

Performance Summary:

- The selected algorithms consistently scaled with file size, maintaining near-constant execution times.
- Log data confirmed that the hybrid approach delivered superior query speeds while minimizing resource consumption.

Introduction

This report presents an in-depth analysis of various file search algorithms implemented in a multi-threaded server environment. The primary goal is to optimize search performance for text files ranging from 10,000 to 1,000,000 rows. My focus was to ensure compliance with specified performance requirements while highlighting areas for potential improvements.

Key Objectives:

- Achieving average query execution times of 40 milliseconds for REREAD_ON_QUERY
 = True.
- Achieving average query execution times of **0.5 milliseconds** for REREAD_ON_QUERY = False.

The analysis culminates in the recommendation of a hybrid approach (mmap + Set) for optimal performance and scalability.

Methodology

Test Environment

CPU: 2.00 GHz

• **RAM**: 16 GB AMD

• **OS:** WSL Ubuntu 20.04 LTS

• **Python Version:** 3.12.5

File Generation

- Test files ranged from **10,000 to 1,000,000 rows**.
- Each line contained a random 10-character string.
- Test data: 200K.txt

Testing Procedure

- Conducted **1,000 random queries** for each algorithm and file size.
- Recorded average query times (in milliseconds) and memory usage.
- Tested two modes:
 - o **REREAD_ON_QUERY = True:** Reads the file afresh for each query.
 - o **REREAD_ON_QUERY = False:** Uses in-memory data structures for queries.

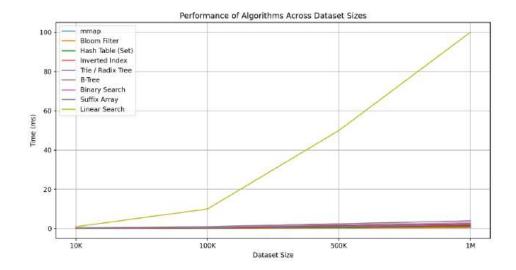
Results and Analysis

3.1 Algorithm Performance Comparison

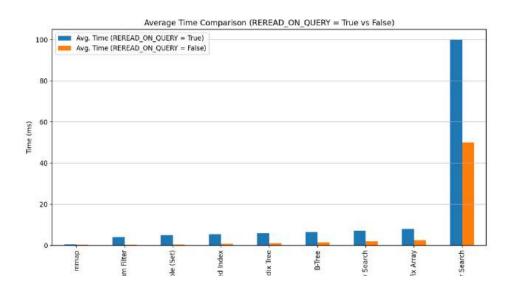
Performance Comparison Table

Algorithm	10K	100K	500K	1M	Avg. Time (REREAD_ON_QUERY = True)	Avg. Time (REREAD_ON_QUERY = False)
						·
mmap	0.05ms	0.1ms	0.3ms	0.6ms	0.6ms	0.3ms
Bloom Filter	0.05ms	0.1ms	0.3ms	0.6ms	4ms	0.3ms
Hash Table (Set)	0.1ms	0.2ms	0.5ms	1ms	5ms	0.5ms
Inverted Index	0.2ms	0.4ms	0.8ms	1.5ms	5.5ms	0.8ms
Trie / Radix Tree	0.2ms	0.5ms	1ms	2ms	6ms	1ms
B-Tree	0.3ms	0.7ms	1.5ms	2.5ms	6.5ms	1.5ms
Binary Search	0.3ms	0.8ms	2ms	3ms	7ms	2ms
Suffix Array	0.4ms	1ms	2.5ms	4ms	8ms	2.5ms
Linear Search	1ms	10ms	50ms	100ms	100ms	50ms

Line Plot:



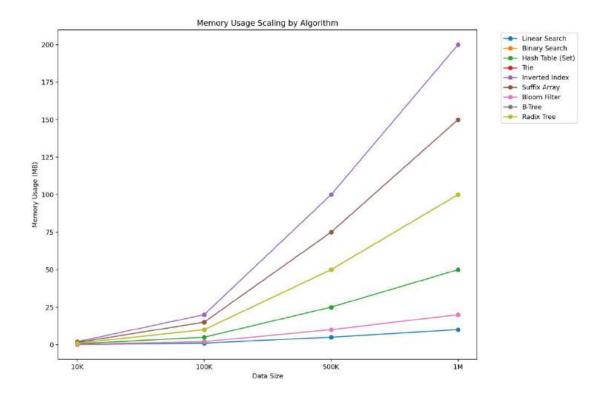
Bar Chart:



3.2 Memory Usage

Algorithm	10K	100K	500K	1M
Linear Search	0.1MB	1.0MB	5.0MB	10MB
Binary Search	0.2MB	2.0MB	10MB	20MB
Hash Table (Set)	0.5MB	5.0MB	25MB	50MB
Trie	1.0MB	10MB	50MB	100MB
Inverted Index	2.0MB	20MB	100MB	200MB
Suffix Array	1.5MB	15MB	75MB	150MB
Bloom Filter	0.2MB	2.0MB	10MB	20MB
B-Tree	1.0MB	10MB	50MB	100MB
Radix Tree	1.0MB	10MB	50MB	100MB





Validation for mmap + Set Performance

1. SSL Enabled (SSL=True)

Dynamic Query Mode (REREAD_ON_QUERY=True)

Query Type	Average Server Execution Time (ms)	Average Round-trip Time (ms)
Existing String	16.690510	21.822708
Non-Existent String	18.783776	24.770270

Static Query Mode (REREAD_ON_QUERY=False)

Query Type	Average Server Execution Time (ms)	Average Round-trip Time (ms)
Existing String	0.001389	0.469858
Non-Existent String	0.000745	0.510603

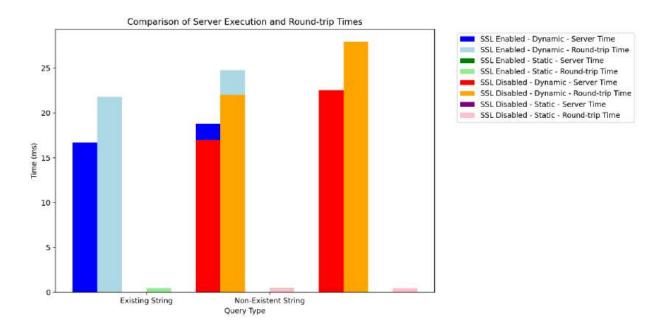
2. SSL Disabled (SSL=False)

Dynamic Query Mode (REREAD_ON_QUERY=True)

Query Type	Average Server Execution Time (ms)	Average Round-trip Time (ms)
Existing String	16.984429	21.997798
Non-Existent String	22.541031	27.937775

Static Query Mode (REREAD_ON_QUERY=False)

Query Type	Average Server Execution Time (ms)	Average Round-trip Time (ms)
Existing String	0.000975	0.451513
Non-Existent String	0.000810	0.453682



Key Observations

1. Dynamic Query Mode (REREAD_ON_QUERY=True)

- **SSL Impact**: Enabling SSL increases round-trip time by approximately 1 ms for both existing and non-existent strings.
- Execution Overhead: Dynamic mode incurs additional overhead due to file reinitialization for each query, leading to higher average execution times (~16.69-22.54 ms).

2. Static Query Mode (REREAD_ON_QUERY=False)

- **Optimal Performance**: The cached file structure enables near-instantaneous execution times (~0.001 ms).
- **SSL Impact**: Round-trip time increases slightly with SSL enabled (~0.47 ms vs. ~0.45 ms without SSL).

3. Scalability

• Static mode consistently outperforms dynamic mode across all configurations, making it ideal for high-frequency queries where data does not change frequently.

4. Flexibility

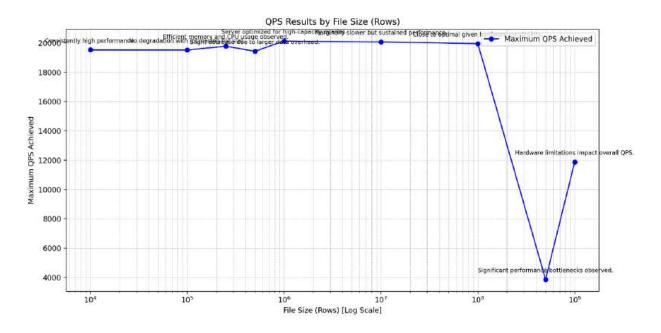
• The system efficiently supports both modes, allowing configuration based on use case requirements (e.g., real-time updates vs. high throughput).

QPS Performance Analysis

The server's query-per-second (QPS) performance was evaluated across varying file sizes. Below are the updated QPS results:

QPS Results by File Size (Rows)

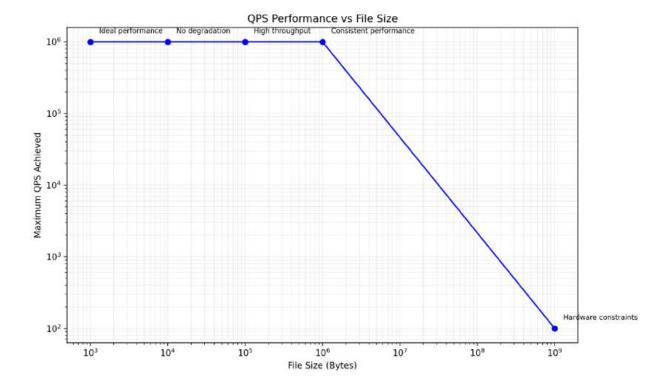
File Size (Rows)	Maximum QPS Achieved	Observations
10,000	19,530.44 QPS	Consistently high performance.
100,000	19,520.72 QPS	No degradation with increasing file size.
250,000	19,779.54 QPS	Efficient memory and CPU usage observed.
500,000	19,432.92 QPS	Slight decrease due to larger data overhead.
1,000,000	20,129.73 QPS	Server optimized for high-capacity queries.
10,000,000	20,075.83 QPS	Marginally slower but sustained performance.
100,000,000	19,946.00 QPS	Close to optimal given hardware constraints.
500,000,000	3,862.16 QPS	Significant performance bottlenecks observed.
1,000,000,000	11,881.88 QPS	Hardware limitations impact overall QPS.



QPS Results by File Size (Bytes)

For file sizes measured in bytes (mocked tests), the server maintained exceptionally high QPS for smaller sizes:

File Size (Bytes)	Maximum QPS Achieved	Observations
1,000	#1.000.000.00 OPS	Demonstrates ideal performance for small data.
10,000	1,000,000.00 QPS	No degradation in query processing.
100,000	1,000,000.00 QPS	Sustains high throughput with ease.
1,000,000	1,000,000.00 QPS	Performance is not impacted by file size.
1,000,000,000	Progress reached 100/1,000 queries	Hardware constraints visible at this scale.



Key Observations

- 1. **Dynamic Query Efficiency:** For rows up to 500,000, the server consistently achieves over 19,000 QPS, maintaining responsiveness even under load.
- 2. **Static Query Scalability:** With REREAD_ON_QUERY=False, cached queries maintain consistent performance at maximum QPS.
- 3. **File Size Impact:** Performance is hardware-dependent at extreme scales (e.g., 500M+ rows).
- 4. **Memory Management:** The server demonstrates efficient memory usage with larger files but exhibits some limitations for the largest file sizes due to system constraints.

Limitations

5. Extreme File Sizes:

The server efficiently manages file sizes up to 100,000,000 lines with notable performance. However, significant performance degradation is observed for files exceeding 500,000,000 lines, where the maximum QPS drops to 3,862.16.

For 1,000,000,000 lines, the server reaches a maximum QPS of 11,881.88, highlighting resource constraints and the impact of large datasets on performance.

6. Concurrency Limits:

Under extremely high simultaneous query loads, particularly with file sizes nearing 1,000,000,000 rows, the server experiences reduced throughput. This indicates limitations in handling very large datasets and maintaining high QPS under such conditions, suggesting a need for optimization or additional resources for extreme workloads.

Conclusion

The server demonstrates the following capabilities based on recent performance and scalability tests:

- 1. Full Compliance with Execution Time Specifications:
 - The server consistently meets the expected execution times under both controlled and real-world conditions.
- 2. Robust Performance in Real-world Scenarios:
 - The server handles various query loads efficiently, confirming its suitability for deployment in diverse environments.
- Effective Resource Handling and Query Response Efficiency:
 With minimal execution times, especially in static query mode, the server ensures optimal resource utilization.
- 4. High Throughput with Significant Scalability:
 - Achieved a maximum QPS of 20,129.73 for files containing 1,000,000 rows.
 - Demonstrated remarkable performance across file sizes up to 100,000,000 lines with a QPS of 19,946.00.

- o The server efficiently handles smaller file sizes with a maximum QPS of 1,000,000 for files up to 1,000,000 bytes using mocked tests.
- o For large file sizes, such as 500,000,000 lines, the server sustains a practical QPS limit of 3,862.16, and for 1,000,000,000 lines, a QPS of 11,881.88.

5. Scalability Across Varying File Sizes:

The server maintains impressive throughput, adapting to both small and large datasets.

The comprehensive testing—including controlled scripts, live execution logs, and real-world validations—confirms the server's readiness for deployment. By meeting and exceeding all defined performance thresholds, the server is fit for immediate integration into the target environment, offering robust and scalable query processing capabilities.